

```
In [9]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
import math
import random
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
In [10]: location="linear-regression.csv"
df = pd.read_csv('linear-regression.csv')
df.insert(0, 'bias', '1')
```

```
In [11]: df.head()
```

```
Out[11]:
```


	bias	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
<b>0</b>	1	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9
<b>1</b>	1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9
<b>2</b>	1	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9
<b>3</b>	1	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9
<b>4</b>	1	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9

```
In [12]: data=np.array(df,dtype=np.float64)
```

```
In [13]: X=data[:,12]
Y=data[:,1]
print(X.shape)
print(Y.shape)
```

```
(1599, 12)
(1599,)
```

```
In [14]: plt.xlabel('x coordinates')
plt.ylabel('quality')
plt.plot(X[:,1:10], Y, 'bo')
plt.show()
```



```
In [15]: X_train, X_temp, y_train, y_temp = train_test_split(X, Y, test_size=0.5, random_state=42)
```

```
In [16]: X_validation, X_test, y_validation, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
```

```
In [17]: print(len(X_train))
          print(len(X_validation))
          print(len(X_test))
```

```
799
480
320
```

```
In [18]: def h(x, theta):
          return np.round(np.matmul(x, theta))
          '''return ans'''
```

```
In [19]: def cost_function(x, y, theta):
          J=((h(x, theta)-y)).T*(h(x, theta)-y))/(2*len(y))
          return J
```

```
In [20]: def accuracy(x, y, theta):
          l=len(x)
          correct=0
          for i in range(0, l):
              y_predicted=round(h(x[i], theta))
              if(y_predicted==y[i]):
                  correct+=1
          percentage=correct/l
          return percentage
```

```
In [21]: def R_Square(x,y,theta):
          y_mean=round(np.mean(y))
          l=len(y)
          numerator=0
          denominator=0
          for i in range(l):
              numerator=numerator+(y[i]-(h(x[i],theta)))**2
              denominator=denominator+(y[i]-y_mean)**2
          value = numerator/denominator
          return 1-value
```

```
In [22]: def RMSE(x,y,theta):
          l=len(y)
          total=0
          for i in range(l):
              y_predicted=round(h(x[i],theta))
              total=total+(y[i]-y_predicted)**2
          total=total/l
          total=math.sqrt(total)
          return total
```

## Analytical Solution

```
In [23]: theta = np.linalg.inv(X_train.T.dot(X_train)).dot(X_train.T).dot(y_train)
          print(theta)
```

```
[ 2.33689602e+01  3.05502970e-02 -1.19741412e+00 -1.87228641e-01
  1.17612142e-02 -1.81714527e+00  6.11854741e-03 -4.06437389e-03
 -1.90187880e+01 -4.42172779e-01  8.79804133e-01  2.64603431e-01]
```

```
In [24]: #validation set R_Square
          print("R_Square score on validation set "+ str(R_Square(X_validation,y_validation,theta)))

          #test set R_Square
          R_Square(X_test,y_test,theta)
          print("R_Square score on test set "+ str(R_Square(X_test,y_test,theta)))
```

```
R_Square score on validation set 0.3252688172043011
R_Square score on test set 0.4125
```

```
In [25]: #validation set RMSE

          print("RMSE score on validation set "+ str(RMSE(X_validation,y_validation,theta)))

          #test set RMSE

          print("RMSE score on test set "+ str(RMSE(X_test,y_test,theta)))
```

```
RMSE score on validation set 0.7231297716638879
RMSE score on test set 0.6637959023675877
```

## With Regularization

```
In [26]: Identity_matrix = np.identity(len(X_train.T.dot(X_train)), dtype = float)
         lamda=2
```

```
In [27]: theta_regularised = np.linalg.inv((X_train.T.dot(X_train))+(lamda*Identity_matrix)).dot(X_train.T.dot(y_train))
```

```
In [28]: #validation set R_square
         print("R_square score on validation set "+str(R_Square(X_validation,y_validation,theta_regularised)))

         #test set R_square
         print("R_square score on test set "+str(R_Square(X_test,y_test,theta_regularised)))
```

R\_square score on validation set 0.31720430107526887  
R\_square score on test set 0.4041666666666667

```
In [29]: #validation set RMSE
         print("RMSE score in validation set "+str(RMSE(X_validation,y_validation,theta_regularised)))

         #test set RMSE
         print("RMSE score in test set "+str(RMSE(X_test,y_test,theta_regularised)))
```

RMSE score in validation set 0.7274384280931732  
RMSE score in test set 0.6684870978560469

## Iterative Gradient Ascent

```
In [30]: data=np.array(df,dtype=np.float64)
         scaler=StandardScaler()
         #data=scaler.fit_transform(data)
         X=data[:,12]
         X=scaler.fit_transform(X)
         Y=data[:,1]
         X[:,0]+=1
         X_train, X_temp, y_train, y_temp = train_test_split(X, Y, test_size=0.5, random_state=42)
         X_validation, X_test, y_validation, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
         #scaler=StandardScaler()
         #X_train=scaler.fit_transform(X_train)
```

```
In [31]: def predict_y(x,theta):
         t=np.dot(x,theta)
         #t=np.round(t)
         return t
```


```
In [32]: def total_cost(y,x,theta):  
    l=len(y)  
    total=0  
    for i in range(l):  
        y_p=predict_y(x[i],theta)  
        total+=((y[i]-y_p)**2)  
    return total/(2*l)
```

```
In [33]: def gradient_ascent(x,y,theta,learning_rate,epoch):  
    l=len(x)  
    cost=[]  
    cost_v=[]  
    y_ep=[]  
    for k in range(epoch):  
        for i in range(0,l):  
            y_p=(np.dot(x[i],theta))  
            theta=theta+learning_rate*((y[i]-y_p)*x[i])  
            cost.append(total_cost(y,x,theta))  
            cost_v.append(total_cost(y_validation,X_validation,theta))  
            y_ep.append(k+1)  
    plt.xlabel('Epochs')  
    plt.ylabel('Cost')  
    plt.plot(y_ep, cost, 'm', linewidth = "5")  
    plt.plot(y_ep, cost_v, 'g', linewidth = "5")  
    plt.show()  
  
    return theta
```

## learning rate 0.01

```
In [34]: theta=np.ones(X_train.shape[1])  
    print(theta)  
    theta=gradient_ascent(X_train,y_train,theta,0.01,30)  
    #theta=gradient_ascent(X_validation,y_validation,theta,0.001,8)  
    print(theta)
```

```
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```



In [35]:

```
#validation set R_square
print("R_square score on validation set "+str(R_Square(X_validation,y_validation,theta)))

#test set R_square
print("R_square score on test set "+str(R_Square(X_test,y_test,theta)))

#validation set RMSE
print("RMSE score in validation set "+str(RMSE(X_validation,y_validation,theta)))

#test set RMSE
print("RMSE score in test set "+str(RMSE(X_test,y_test,theta)))
```


```
R_square score on validation set 0.303763440860215
R_square score on test set 0.3666666666666667
RMSE score in validation set 0.734563362367967
RMSE score in test set 0.689202437604511
```

## learning rate 0.001

In [36]:

```
theta=np.zeros(X_train.shape[1])
print(theta)
theta=gradient_ascent(X_train,y_train,theta,0.001,25)
#theta=gradient_ascent(X_validation,y_validation,theta,0.001,8)
print(theta)
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```



```
[ 5.63609973  0.08777612 -0.22136512 -0.03556634  0.01875345 -0.08741357
 0.05674787 -0.13166918 -0.06067875 -0.06059899  0.15495999  0.2621672 ]
```

In [37]:

```

#validation set R_square
print("R_square score on validation set "+str(R_Square(X_validation,y_validation,theta)))

#test set R_square
print("R_square score on test set "+str(R_Square(X_test,y_test,theta)))

#validation set RMSE
print("RMSE score in validation set "+str(RMSE(X_validation,y_validation,theta)))

#test set RMSE
print("RMSE score in test set "+str(RMSE(X_test,y_test,theta)))

```

R\_square score on validation set 0.3360215053763441

R\_square score on test set 0.4083333333333333

RMSE score in validation set 0.7173446405552447

RMSE score in test set 0.6661456297237114

## learning rate 0.0001


In [38]:

```

theta=np.zeros(X_train.shape[1])
print(theta)
theta=gradient_ascent(X_train,y_train,theta,0.0001,100)
#theta=gradient_ascent(X_validation,y_validation,theta,0.001,8)
print(theta)

```

[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]



```

[ 5.6322383  0.10694348 -0.20932108 -0.03079476  0.02918622 -0.08333899
 0.05566896 -0.12365031 -0.0847903  -0.0332272  0.15562779  0.25233705]

```

In [39]:

```
#validation set R_square
print("R_square score on validation set "+str(R_Square(X_validation,y_validation,theta)))

#test set R_square
print("R_square score on test set "+str(R_Square(X_test,y_test,theta)))

#validation set RMSE
print("RMSE score in validation set "+str(RMSE(X_validation,y_validation,theta)))

#test set RMSE
print("RMSE score in test set "+str(RMSE(X_test,y_test,theta)))
```

R\_square score on validation set 0.3198924731182796

R\_square score on test set 0.4041666666666667

RMSE score in validation set 0.7260050504874834

RMSE score in test set 0.6684870978560469

In [ ]:

In [ ]: